

Large-Scale Machine Learning

What do we need to do differently when $m \sim \Theta(10^8)$?

- Using a small data set chosen from the larger set, plot training curves and verify that the model has high variance.
If it does not, a randomly-chosen subset will probably work

Stochastic Gradient Descent (as opposed to "Batch" gradient descent)

If m is large, fitting takes a long time.

$$\text{cost}(\theta, (x^{(i)}, y^{(i)})) = \frac{1}{2} (h_\theta(x^{(i)}) - y^{(i)})^2 \rightarrow J_{\text{TRAIN}} = \frac{1}{m} \sum_{i=1}^m \text{cost}(\theta, (x^{(i)}, y^{(i)}))$$

1. randomly shuffle dataset

2. Repeat { % 1-10 passes common

for $i = 1:m$

$$\theta = \theta - \alpha (h_\theta(x^{(i)}) - y^{(i)}) x^{(i)}$$

{ % take a step for each data point}

The system converges to a random walk in the region of the global minimum.

Mini-Batch Gradient Descent





- Batch - use all m for each iteration
- Stochastic - use 1 for each iteration
- Mini-Batch - Use b examples in each loop ($b \in [2, 100]$) and use a new subset for each iteration.

* Tends to out-perform Stochastic gradient descent when the computation can be parallelized. *

Stochastic Gradient Descent - Convergence

Computing J_{TRAIN} is expensive.

During learning, compute cost before update. Plot cost averaged over a number of steps.



Slowly decreasing α over time can improve convergence. (not common)

$$\text{E.g. } \alpha = \frac{\text{const 1}}{\text{iteration\#} + \text{const 2}}$$

Online Learning - continuous stream of data

Repeat forever {

- Get (x, y)
 - Update θ using (x, y)
- }
- don't save training data
 - θ evolves over time (similar to stochastic gradient descent)
 - * can adapt to changing user preferences *

CTR → "Click Through Rate"

Example: Product Search (learning to search)

- User searches for a product
- Return top results
- make a feature vector for each returned result
 - product characteristics
 - count of words in description matching search terms
 - etc
- $y = 1$ if user clicks on link, 0 otherwise
- learn $p(y=1|x; \theta)$
- Use θ to improve searches, or choose special offers etc.

May be useful in conjunction with collaborative learning.

Map Reduce and Data Parallelism (Jeffrey Dean and Sanjay Ghemawat)

* Some machine learning algorithms are too large for one machine. *

Iterate!

- Divide data evenly among available machines/cores.
- On each machine:
 - iterate using its subset of data
 - send result for subset to master node
- master node combines result and updates parameters

* HADOOP *

Many learning algorithms can be expressed as a sum of functions over the training set. These are candidates for map-reduce parallelization.

Example: Neural Network

- do forward/back propagation on subsets of data
- combine results on master node